



London
Business
School

AM05 Data Management

02. Logical Database Design

Dr. David Tilson

London
Business
School

AM05: Data Management


David Tilson

Welcome!


Class starts at 12:45 (London). Thank you for being early!

Zoom Classroom Etiquette

1. Roomies: Please **turn on** your cameras and **join without audio**
2. Rename yourself
 - a. Add "R" (for Roomie) in front of your name if you're in the LT
 - b. Add "Z" (for Zoomie) in front of your name if you're remote
3. Questions
 - a. Raise you (digital) hand if you want to speak or ask a question that needs an answer ASAP
 - b. Use chat to ask questions that can wait for a few minutes
 - c. I will ask you to answer questions using options on "participants" panel
 - d. Technical issues - message the facilitator privately in Zoom chat



Roomies
When prompted, click "X"



 Raise Hand



4. If you are in a breakout room, engage with your colleagues to extract the most out the class
5. Session will be recorded

¶ Components of Relational Data Model

- **Data structure** – Relations (tables), rows, and columns
- **Data manipulation** – SQL operations for retrieving/modifying data
- **Data integrity** – Mechanisms for implementing business rules that maintain integrity of manipulated data

¶ Converting ERD into Database

¶ Database Normalization

¶ Hands On

¶ For next time

73

In the relational model data is represented in **relations (or tables)** which have certain properties

1. Each **relation (table)** has a unique name

2. Each **attribute (column)** has a unique name

3. An entry at an intersection of row and column is **atomic** (not multivalued).

Employee

| | EmpNo | EmpName | Department | Email | Phone |
|---|-------|--------------------|----------------|--------------------|--------|
| ▶ | E100 | Chuck Coordinator | Administration | chuck@colorado.edu | 3-1111 |
| | E101 | Mary Manager | Football | mary@colorado.edu | 5-1111 |
| | E102 | Sally Supervisor | Planning | sally@colorado.edu | 3-2222 |
| | E103 | Alan Administrator | Administration | alan@colorado.edu | 3-3333 |

4. Each **row (record / tuple)** is unique

- Implies that relation must have a **Primary Key (PK)** – an attribute (or combination of attributes) that *uniquely identifies* each row.
- Corresponds to the identifier in an ERD

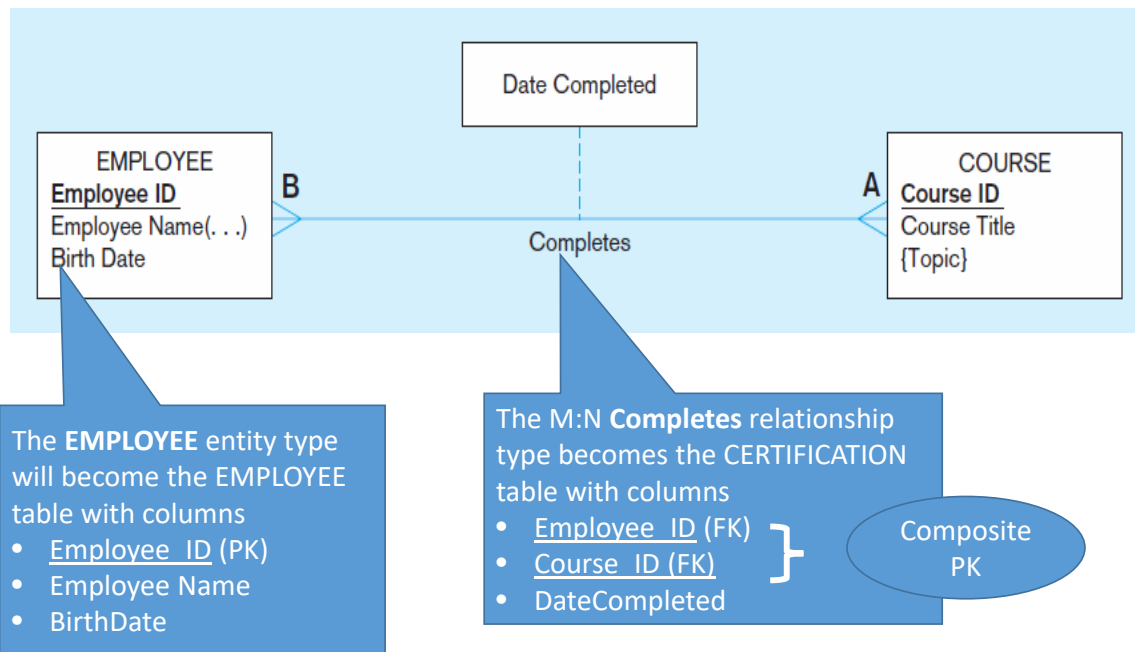
5. Sequence is NOT significant for

- Columns (left-to-right)
- Rows (top-to-bottom)



- All relations are tables. . . but not all tables are relations
- Relations can be thought of as **sets** of tuples
- Relations that follow these rules are in 1st Normal Form (will see soon)

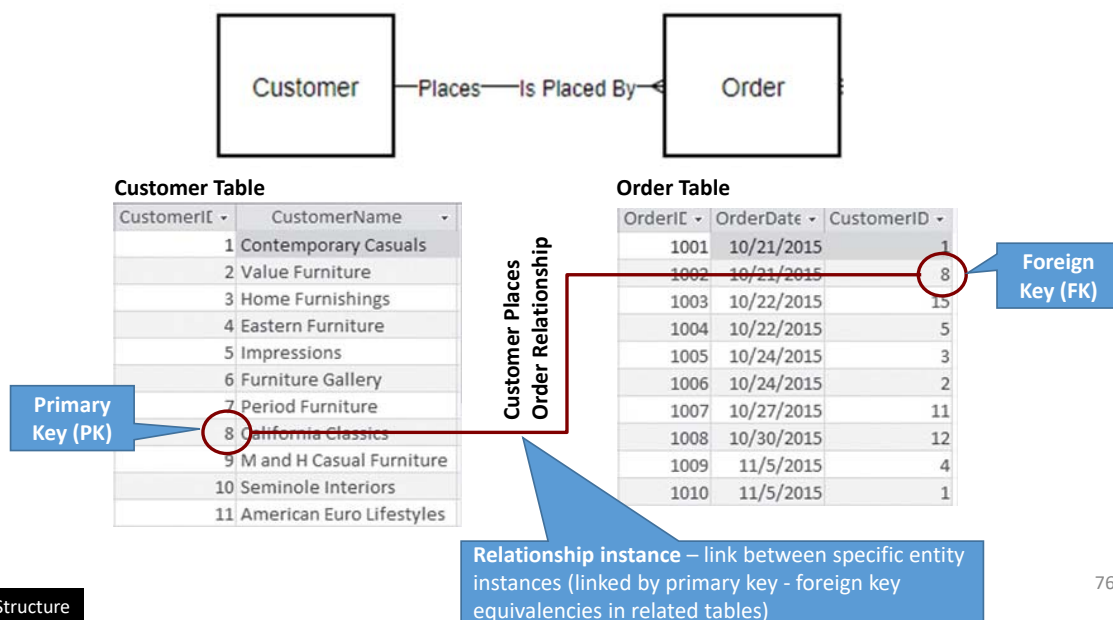
Relations (tables) correspond with **entity types** and with **many-to-many relationship types** (or associative entities)



Data Structure NOTE: The word **relation** (in a relational database) is NOT the same as the word **relationship** in an ERD

75

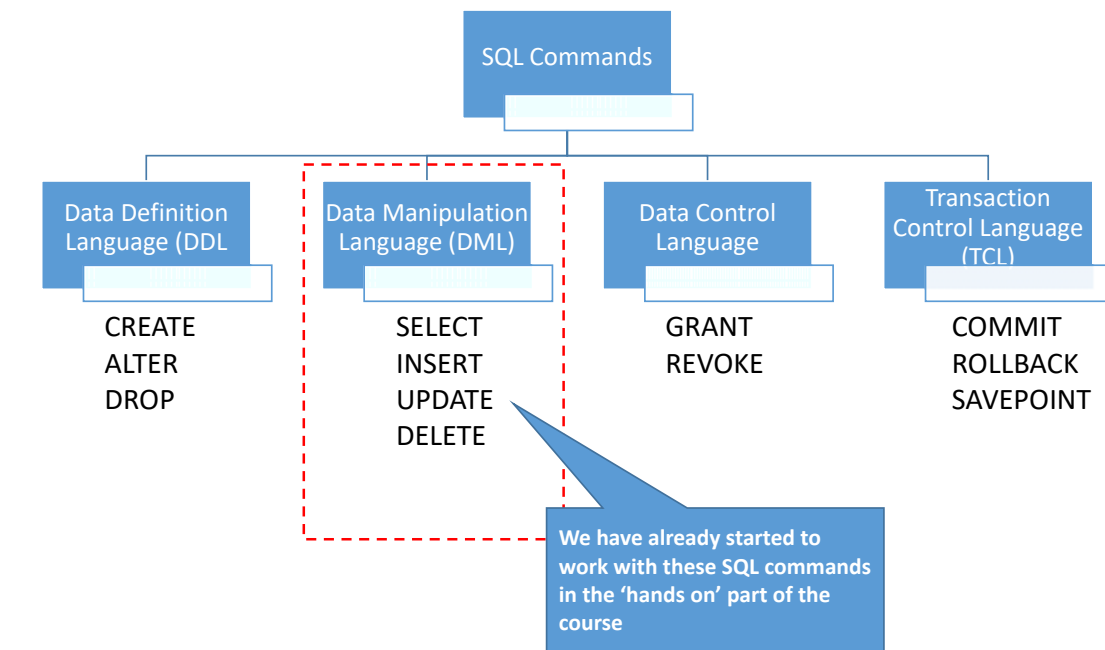
Foreign keys are identifiers that enable a *dependent relation* (on the many side of a relationship) to refer to its *parent relation* (on the one side of the relationship)



Data Structure

76

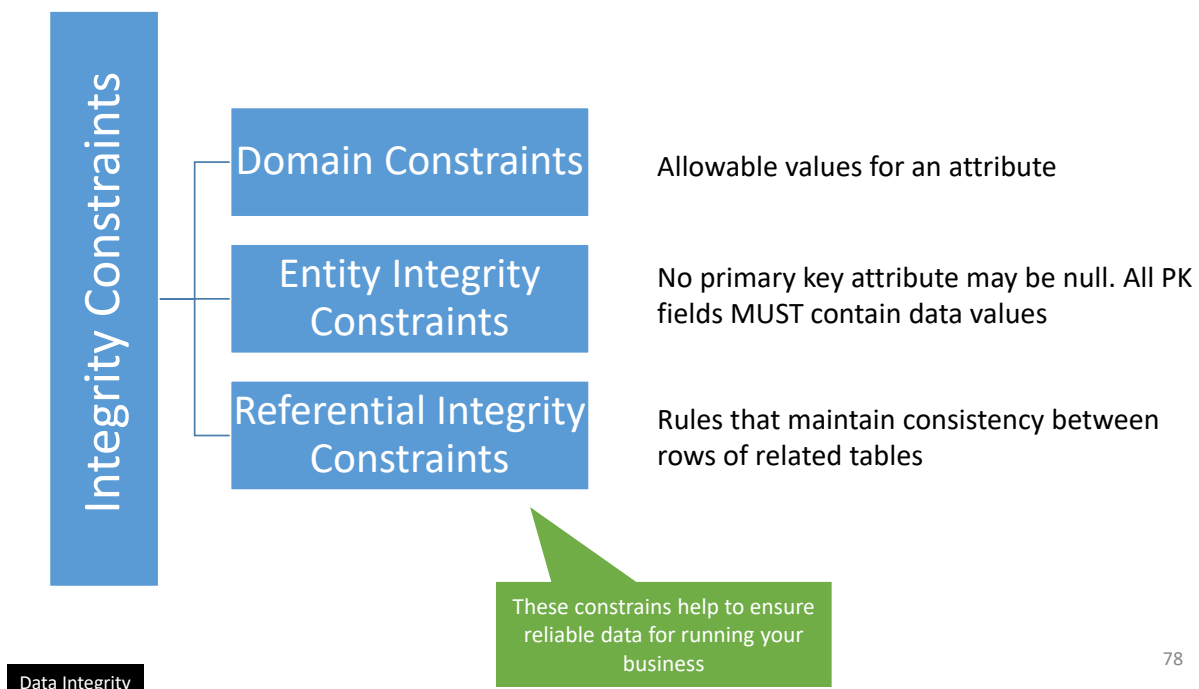
SQL provides powerful operations for retrieving and modifying data



Data Manipulation

77

Relational model includes **constraints** (rules) limiting acceptable values and actions to help **maintain accuracy and integrity** of the data in a database



Data Integrity

78

Data Domain refers to all values a data element may contain.
Rule can be as simple as data type or enumerated list of values

TABLE 4-1 Domain Definitions for INVOICE Attributes

| Attribute | Domain Name | Description | Domain |
|----------------------|----------------------|--|-----------------------|
| CustomerID | Customer IDs | Set of all possible customer IDs | character: size 5 |
| CustomerName | Customer Names | Set of all possible customer names | character: size 25 |
| CustomerAddress | Customer Addresses | Set of all possible customer addresses | character: size 30 |
| CustomerCity | Cities | Set of all possible cities | character: size 20 |
| CustomerState | States | Set of all possible states | character: size 2 |
| CustomerPostalCode | Postal Codes | Set of all possible postal zip codes | character: size 10 |
| OrderID | Order IDs | Set of all possible order IDs | character: size 5 |
| OrderDate | Order Dates | Set of all possible order dates | date: format mm/dd/yy |
| ProductID | Product IDs | Set of all possible product IDs | character: size 5 |
| ProductDescription | Product Descriptions | Set of all possible product descriptions | character: size 25 |
| ProductFinish | Product Finishes | Set of all possible product finishes | character: size 15 |
| ProductStandardPrice | Unit Prices | Set of all possible unit prices | monetary: 6 digits |
| ProductLineID | Product Line IDs | Set of all possible product line IDs | integer: 3 digits |
| OrderedQuantity | Quantities | Set of all possible ordered quantities | integer: 3 digits |

Domain definitions help ensure that data in a database is valid

Enforced by

- Specifying datatypes and length when creating tables
- Enumerating values in another relation (e.g. list of valid US state abbreviations)
- Checks in applications (e.g. ensure that prices and quantities are positive)

Entity integrity rule is to ensure that relations have a primary key that is unique and cannot contain NULLs

Example of how to create a relation in SQL DDL

```
CREATE TABLE Customer_T
  (CustomerID          NUMBER(11,0) NOT NULL,
   CustomerName        VARCHAR2(25) NOT NULL,
   CustomerAddress      VARCHAR2(30),
   CustomerCity         VARCHAR2(20),
   CustomerState        CHAR(2),
   CustomerPostalCode   VARCHAR2(9),
   CONSTRAINT Customer_PK PRIMARY KEY (CustomerID));
```

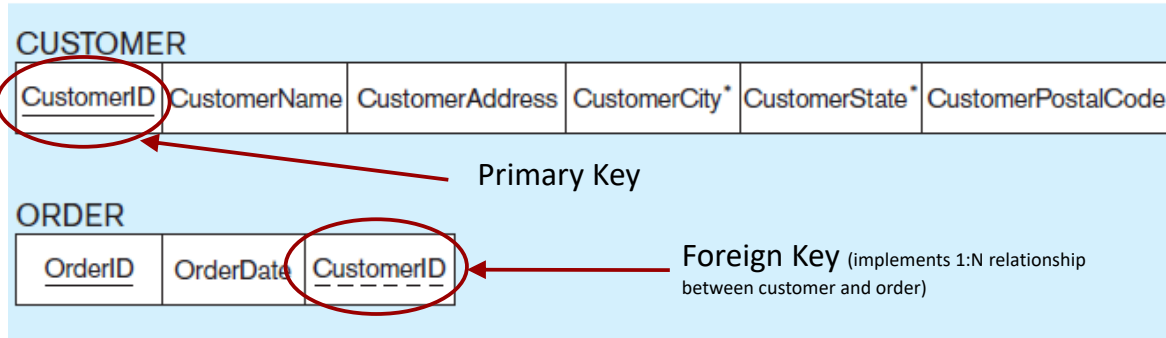
Specify that the CustomerID cannot be NULL

This is one way a Primary Key is specified in SQL

Enforced by

- DBMS not allowing operations (INSERT, UPDATE) to produce an invalid primary key
- Any operation that creates a duplicate primary key or one containing nulls is rejected

Referential integrity rules maintain consistency among the rows of relations



Enforced by

- DBMS not allowing insertion of new row in a table with a Foreign Key (FK) unless the value of the FK matches a Primary Key (PK) in the other table. For example, we would not be allowed to insert a new ORDER unless the CustomerID existed in the CUSTOMER table
- Handling issues with deletion of rows in parent table (e.g. CUSTOMER above)*
 - **Restrict**—don't allow delete of *parent* side if related rows exist in *dependent* side
 - **Cascade**—automatically delete *dependent* side rows corresponding to *parent* side row to be deleted
 - **Set-to-Null**—set FK on dependent side to null if deleting from the parent side (but not allowed for weak entities)

Data Integrity – Referential Integrity Constraint

See: <https://www.mysqltutorial.org/mysql-foreign-key/> for more information

Referential integrity constraints implemented with foreign key to primary key references

```
CREATE TABLE Customer_T
(
    CustomerID          NUMBER(11,0)    NOT NULL,
    CustomerName        VARCHAR2(25)    NOT NULL,
    CustomerAddress      VARCHAR2(30),
    CustomerCity         VARCHAR2(20),
    CustomerState        CHAR(2),
    CustomerPostalCode   VARCHAR2(9),
    CONSTRAINT Customer_PK PRIMARY KEY (CustomerID);

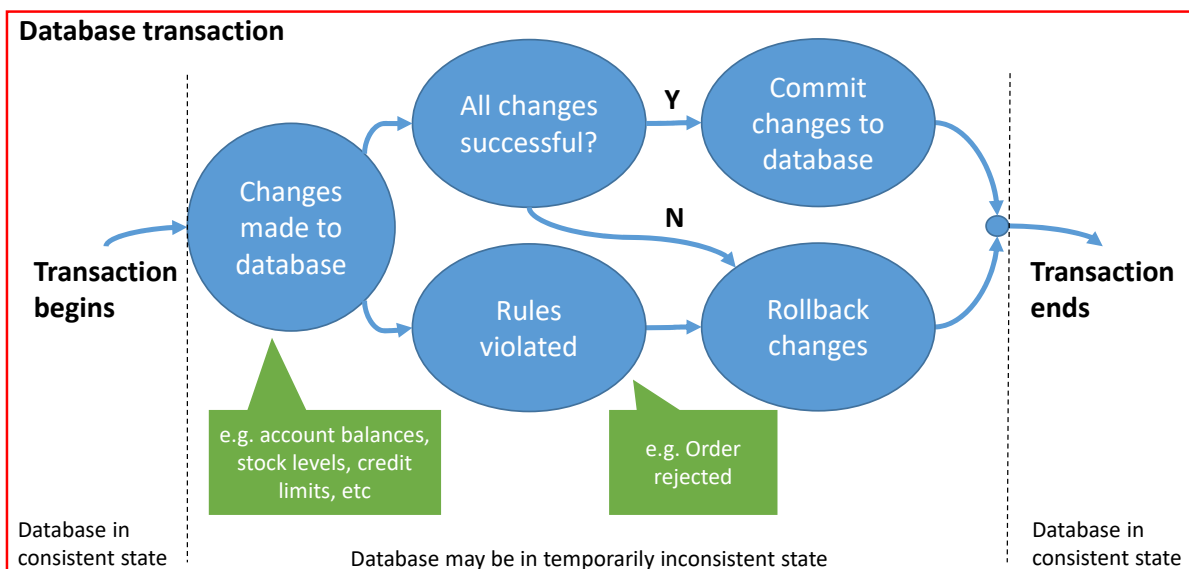
CREATE TABLE Order_T
(
    OrderID             NUMBER(11,0)    NOT NULL,
    OrderDate            DATE DEFAULT SYSDATE,
    CustomerID           NUMBER(11,0),
    CONSTRAINT Order_PK PRIMARY KEY (OrderID),
    CONSTRAINT Order_FK FOREIGN KEY (CustomerID) REFERENCES Customer_T (CustomerID);
```

Quick questions Answer via Zoom chat

1. Why is it important to have “referential integrity constraints” for the relationship between an order and a customer?
2. Would it be a good idea to use VARCHAR as the data type for OrderQuantity? Why/why not?

83

RDMBS can ensure database consistency across several operations using **transactions**




ACID Properties – to ensure accuracy, completeness, and data integrity

| | |
|-------------|---|
| Atomicity | Each transaction is all or nothing |
| Consistency | Data should be consistent according to all rules |
| Isolation | Transactions do not affect each other |
| Durability | Committed data not lost (even if there was a power failure) |

84

E.F. Codd invented the theoretical basis for the relational model



Edgar Frank "Ted" Codd (1923–2003) was an English computer scientist who, while working for IBM, invented the relational model for database management, the theoretical basis for relational databases and relational database management systems.

¶ His 1970 paper **“A Relational Model of Data for Large Shared Data Banks”** introduced the concepts behind relational DBMSs.

¶ He also coined the term OLAP (will discuss this later)

The relational model

- Has its own mathematical underpinning (Relational Algebra)
- Is the basis of most databases in use
- Ensures accurate, consistent data in innumerable mission critical systems

https://en.wikipedia.org/wiki/File:Edgar_F_Codd.jpg

85

¶ Components of Relational Data Model

¶ Converting ERD into Database

¶ Database Normalization


¶ Hands On

¶ For next time

86

Mapping a regular entity to a relation is very simple

CUSTOMER entity type with simple attributes



- Each attribute of the entity becomes a column (field) of the resulting relation
- Identifier of the entity becomes a primary key (PK) in the relation

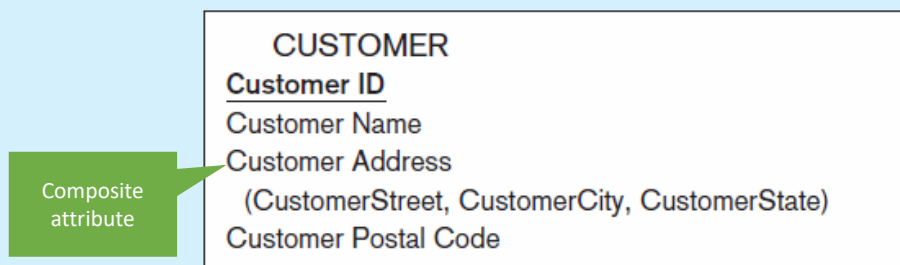
CUSTOMER relation

| CUSTOMER | | | |
|-------------------|--------------|-----------------|--------------------|
| <u>CustomerID</u> | CustomerName | CustomerAddress | CustomerPostalCode |

87

Dealing with a composite attribute is also easy

CUSTOMER entity type with composite attribute



Components of composite attribute become individual columns

CUSTOMER relation with address detail

| CUSTOMER | | | | | |
|-------------------|--------------|----------------|--------------|---------------|--------------------|
| <u>CustomerID</u> | CustomerName | CustomerStreet | CustomerCity | CustomerState | CustomerPostalCode |

88

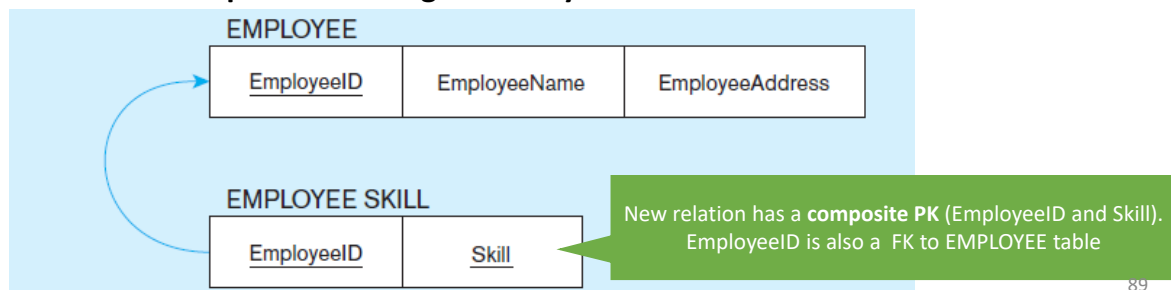
Mapping entity with multivalued attribute achieved with an additional relation

Entity with a multivalued attribute



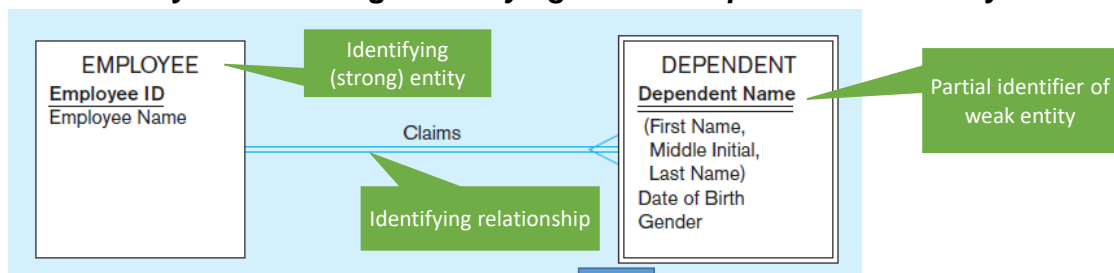
Multivalued attribute becomes a new relation with a 1:M relationship with the original entity

1:M relationship between original entity and new relation

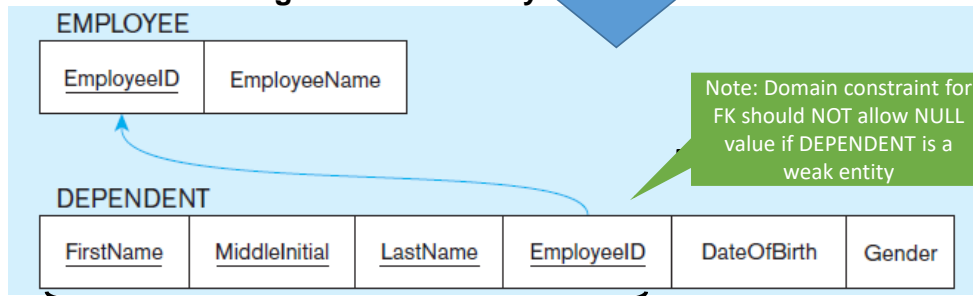


Weak entity needs a composite key with PK of owner entity

Weak entity exists through *identifying relationship* with owner entity



Relations resulting from weak entity

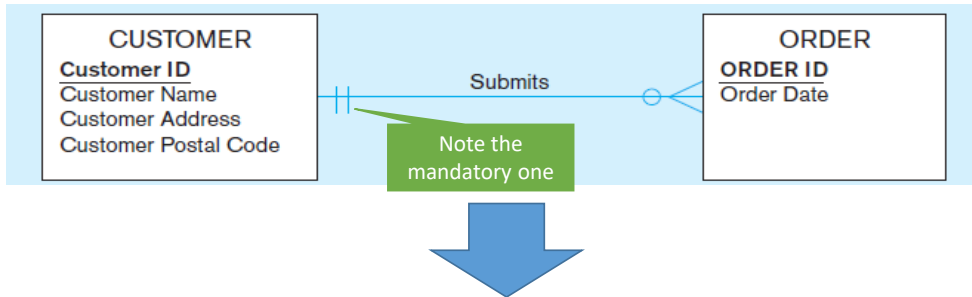


Composite PK composed of

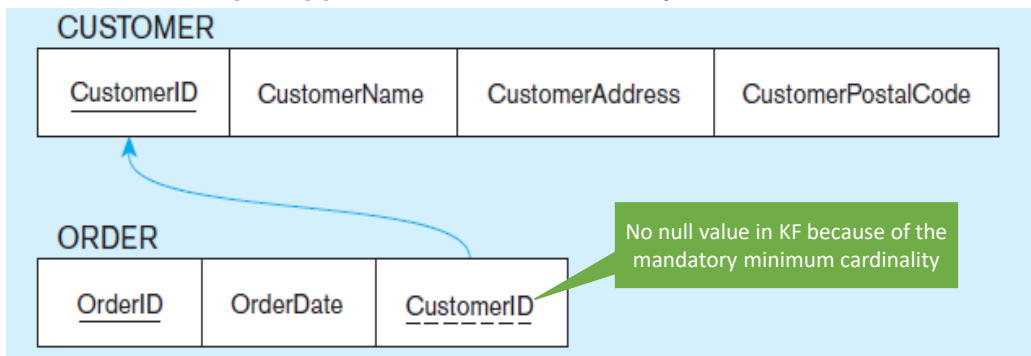
- Partial identifier of weak entity
- PK of identifying relation (strong entity)

For 1:M relationship PK on the *one side* becomes a FK on the *many side*

1:M Relationship between customers and orders



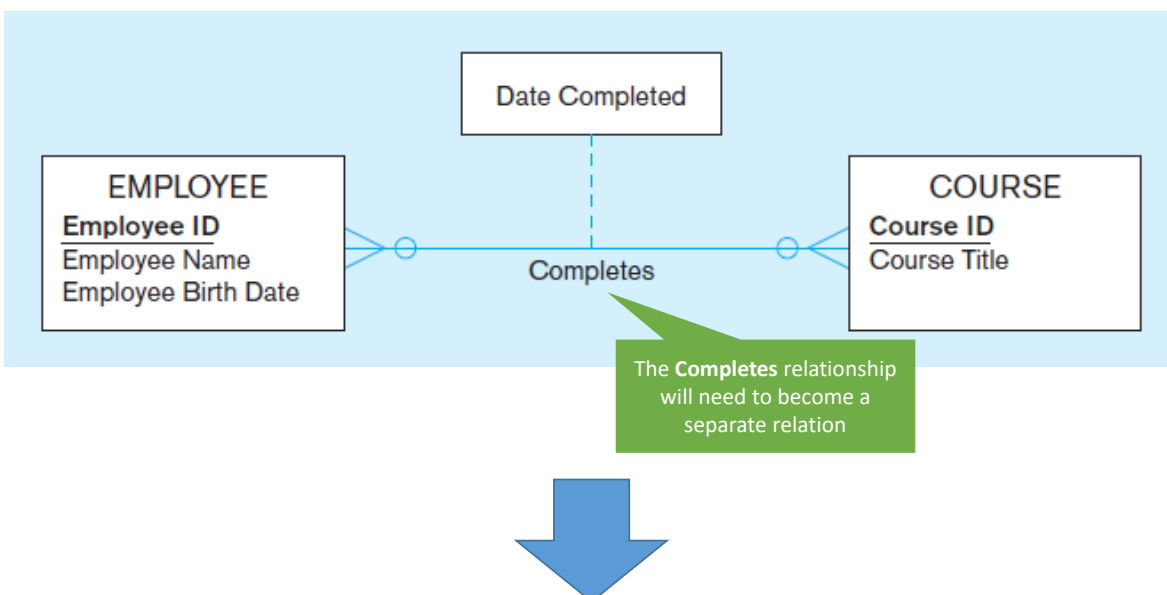
1:M Relationship mapped with FK on the many side



91

M:N relations results in the creation of a new relation with the PKs of the two entities as its PK (1/2)

Completes relationship (M:N)

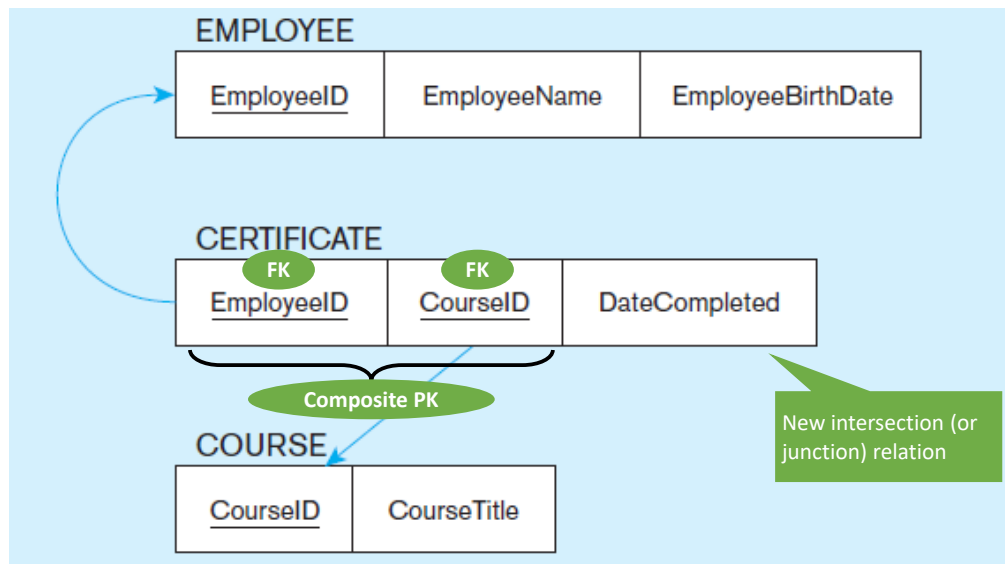


92

M:N relations results in the creation of a new relation with the PKs of the two entities as its PK (2/2)




Three resulting relations



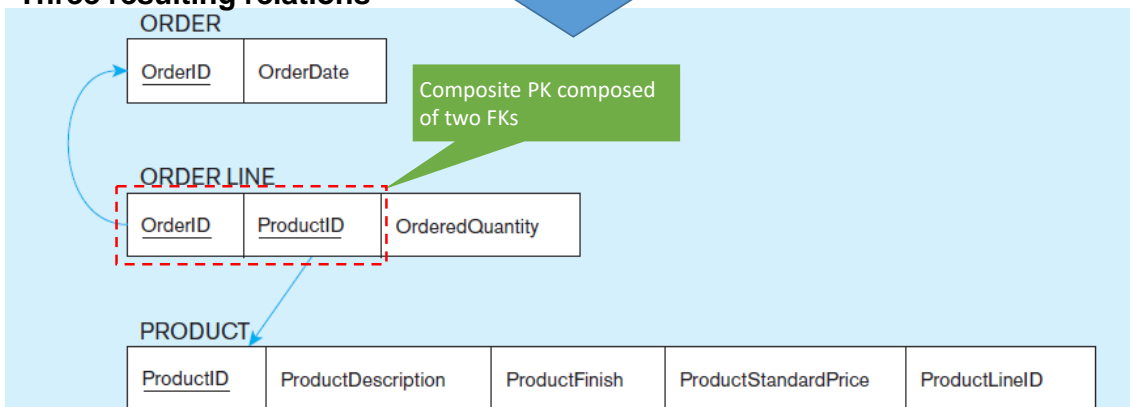
93

Mapping **associative entities** to relational databases is often identical to mapping M:N relationships

Relationship with associative entity



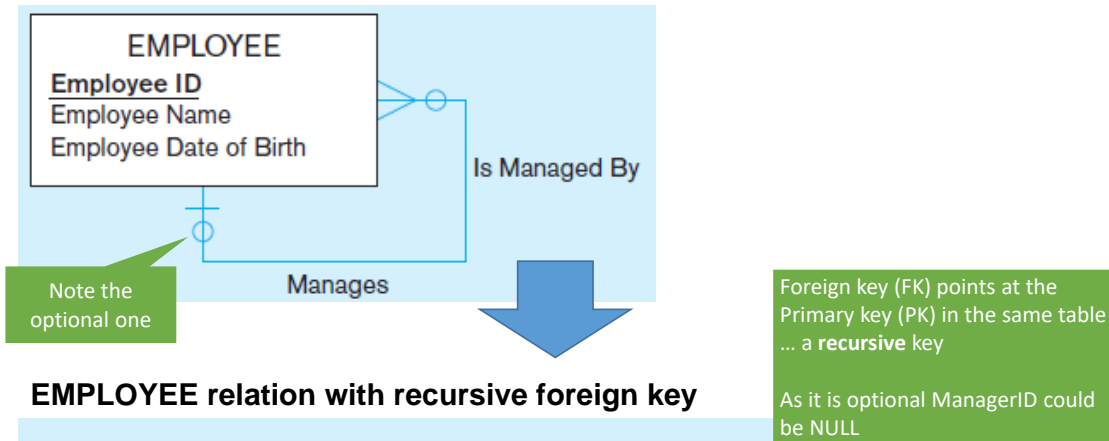
Three resulting relations




94

Unary 1:N relationship can be used to store a hierarchy... in this case supervisor has subordinates, and they in turn could manage others

EMPLOYEE entity with unary 1:N relationship



EMPLOYEE relation with recursive foreign key

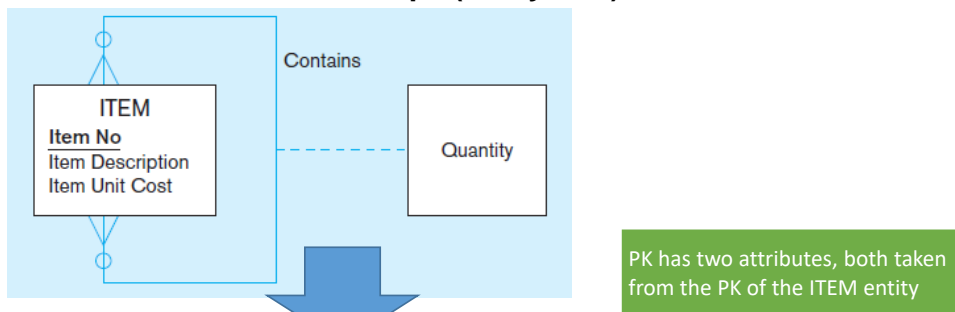


95

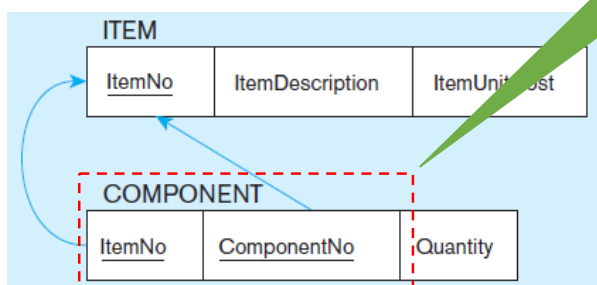
Unary M:N relationship requires two relations

- One for the entity
- One for the associative relation

Bill-of-materials relationships (unary M:N)



ITEM and COMPONENT relations



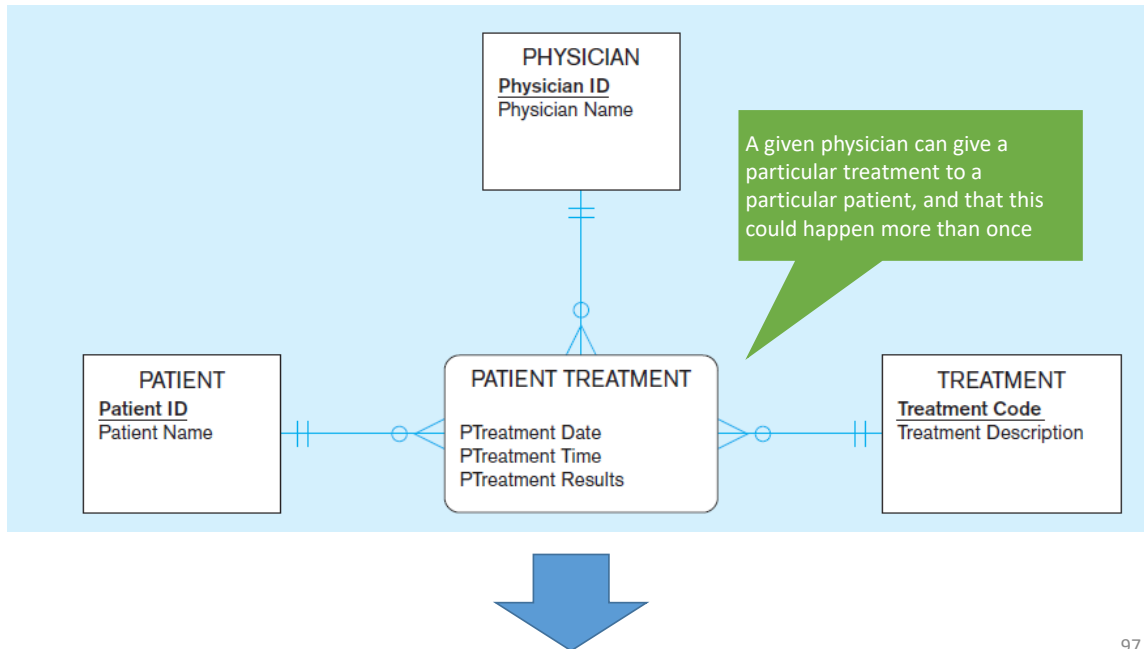
Also example of representing hierarchy

- An item may have sub-items, which could in turn have other sub-items, etc.
- Different from employees example
- Given employee can have only one direct manager (1:N)
- A component could be part of many other items (M:N)

96

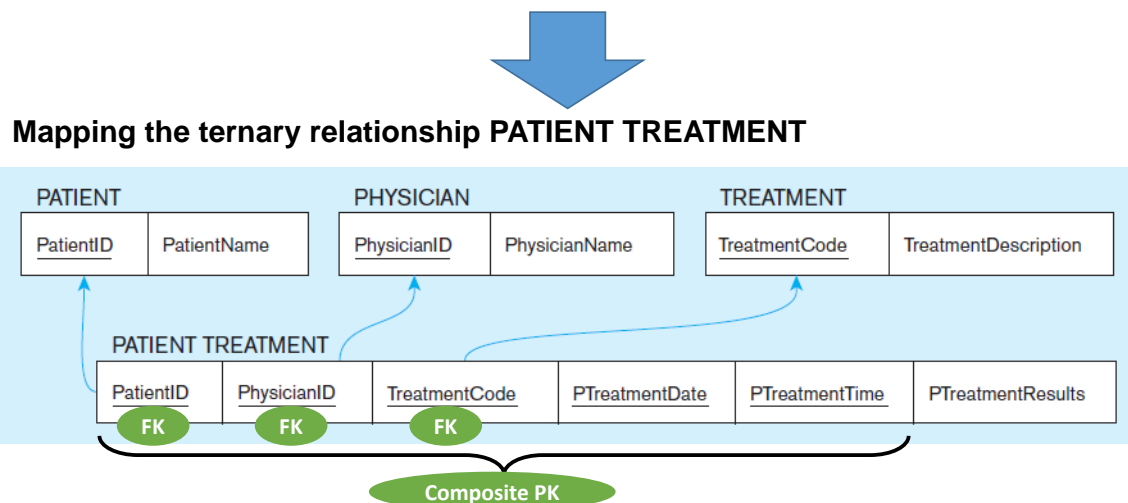
Mapping Ternary (and n-ary) Relationships requires one relation for each entity and one for the associative entity (1/2)

PATIENT TREATMENT Ternary relationship with associative entity



97

Mapping Ternary (and n-ary) Relationships requires one relation for each entity and one for the associative entity (2/2)



Remember that the **PK MUST be unique**... this is why treatment date and time are included in the composite PK

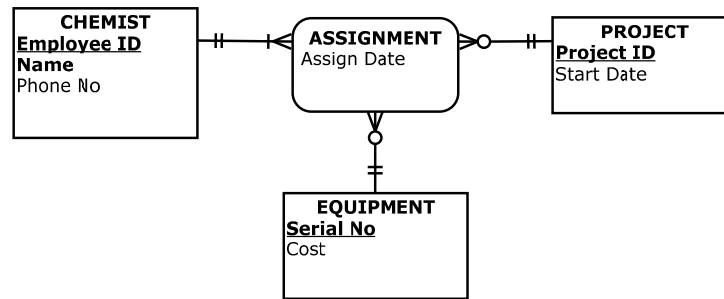
But this makes a very cumbersome key... it would be better to create a surrogate key like TreatmentID

98

In-class exercise: ERD to Schema

10-minute
breakout

- ¶ The following ERD represents a data model for tracking the allocation of laboratory equipment to chemists working on projects.




- ¶ Convert the ERD into a set of relational schemas

Use this format (italics to show FKs)

ORDER OrderID OrderDate *CustomerID*

Which is equivalent to



- ¶ Capture your answer on a PowerPoint page or Document
- ¶ Submit to <https://forms.gle/ts1nwV3m58WXRqRX9> One submission per breakout team

99

¶ Components of Relational Data Model

¶ Converting ERD into Databases

¶ Database Normalization

¶ Hands On

¶ For next time

100

This relation captures employee attributes along with their training history

| EMPLOYEE2 | | | | | |
|---------------|------------------|--------------|--------|--------------------|---------------|
| <u>EmplID</u> | Name | DeptName | Salary | <u>CourseTitle</u> | DateCompleted |
| 100 | Margaret Simpson | Marketing | 48,000 | SPSS | 6/19/2015 |
| 100 | Margaret Simpson | Marketing | 48,000 | Surveys | 10/7/2015 |
| 140 | Alan Beeton | Accounting | 52,000 | Tax Acc | 12/8/2015 |
| 110 | Chris Lucero | Info Systems | 43,000 | Visual Basic | 1/12/2015 |
| 110 | Chris Lucero | Info Systems | 43,000 | C++ | 4/22/2015 |
| 190 | Lorenzo Davis | Finance | 55,000 | | |
| 150 | Susan Martin | Marketing | 42,000 | SPSS | 6/19/2015 |
| 150 | Susan Martin | Marketing | 42,000 | Java | 8/12/2015 |

Is this a relation?



Yes: Unique rows and no multivalued attributes

What's the primary key?



Composite: EmplID, CourseTitle

101

This relation has the potential of suffering from **insertion anomalies**

| EMPLOYEE2 | | | | | |
|---------------|------------------|--------------|--------|--------------------|---------------|
| <u>EmplID</u> | Name | DeptName | Salary | <u>CourseTitle</u> | DateCompleted |
| 100 | Margaret Simpson | Marketing | 48,000 | SPSS | 6/19/2015 |
| 100 | Margaret Simpson | Marketing | 48,000 | Surveys | 10/7/2015 |
| 140 | Alan Beeton | Accounting | 52,000 | Tax Acc | 12/8/2015 |
| 110 | Chris Lucero | Info Systems | 43,000 | Visual Basic | 1/12/2015 |
| 110 | Chris Lucero | Info Systems | 43,000 | C++ | 4/22/2015 |
| 190 | Lorenzo Davis | Finance | 55,000 | | |
| 150 | Susan Martin | Marketing | 42,000 | SPSS | 6/19/2015 |
| 150 | Susan Martin | Marketing | 42,000 | Java | 8/12/2015 |

Insertion anomalies since we cannot insert information on

- A new employee without having an employee taking a class (otherwise CourseTitle would be NULL)
- A new course without an employee taking it (otherwise EmplID would be NULL)

102

It can also suffer from deletion anomalies . . .

| EMPLOYEE2 | | | | | |
|---------------|------------------|--------------|--------|--------------------|---------------|
| <u>EmplID</u> | Name | DeptName | Salary | <u>CourseTitle</u> | DateCompleted |
| 100 | Margaret Simpson | Marketing | 48,000 | SPSS | 6/19/2015 |
| 100 | Margaret Simpson | Marketing | 48,000 | Surveys | 10/7/2015 |
| 140 | Alan Beeton | Accounting | 52,000 | Tax Acc | 12/8/2015 |
| 110 | Chris Lucero | Info Systems | 43,000 | Visual Basic | 1/12/2015 |
| 110 | Chris Lucero | Info Systems | 43,000 | C++ | 4/22/2015 |
| 190 | Lorenzo Davis | Finance | 55,000 | | |
| 150 | Susan Martin | Marketing | 42,000 | SPSS | 6/19/2015 |
| 150 | Susan Martin | Marketing | 42,000 | Java | 8/12/2015 |

Deletion anomaly if we remove employee 140, we lose information about the existence of a Tax Acc class

103

. . . as well as modification anomalies

| EMPLOYEE2 | | | | | |
|---------------|------------------|--------------|--------|--------------------|---------------|
| <u>EmplID</u> | Name | DeptName | Salary | <u>CourseTitle</u> | DateCompleted |
| 100 | Margaret Simpson | Marketing | 48,000 | SPSS | 6/19/2015 |
| 100 | Margaret Simpson | Marketing | 48,000 | Surveys | 10/7/2015 |
| 140 | Alan Beeton | Accounting | 52,000 | Tax Acc | 12/8/2015 |
| 110 | Chris Lucero | Info Systems | 43,000 | Visual Basic | 1/12/2015 |
| 110 | Chris Lucero | Info Systems | 43,000 | C++ | 4/22/2015 |
| 190 | Lorenzo Davis | Finance | 55,000 | | |
| 150 | Susan Martin | Marketing | 42,000 | SPSS | 6/19/2015 |
| 150 | Susan Martin | Marketing | 42,000 | Java | 8/12/2015 |

Modification anomaly since giving a salary increase to employee 100 forces us to update multiple records i.e. there are opportunities for data inconsistencies

104

We get the anomalies because there are two themes (entity types) in this relation. This results in **data duplication** and an unnecessary dependency between the entities

While EMPLOYEE2 is a relation . . . But it is **NOT** a Well-Structured Relation

| EmplID | Name | DeptName | Salary | CourseTitle | DateCompleted |
|--------|------------------|--------------|--------|--------------|---------------|
| 100 | Margaret Simpson | Marketing | 48,000 | SPSS | 6/19/2015 |
| 100 | Margaret Simpson | Marketing | 48,000 | Surveys | 10/7/2015 |
| 140 | Alan Beeton | Accounting | 52,000 | Tax Acc | 12/8/2015 |
| 110 | Chris Lucero | Info Systems | 43,000 | Visual Basic | 1/12/2015 |
| 110 | Chris Lucero | Info Systems | 43,000 | C++ | 4/22/2015 |
| 190 | Lorenzo Davis | Finance | 55,000 | | |
| 150 | Susan Martin | Marketing | 42,000 | SPSS | 6/19/2015 |
| 150 | Susan Martin | Marketing | 42,000 | Java | 8/12/2015 |

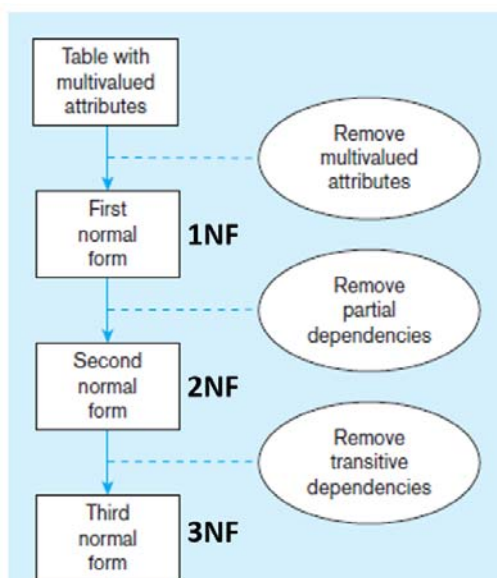
Fields pertaining to employees

Fields pertaining to courses

➡ Rule of thumb: A table should not pertain to more than one entity type

105

Data Normalization improves a logical design to *avoid unnecessary duplication of data*



3NF generally considered to be sufficiently well-structured relations

A **Well-Structured Relation** is one that contains minimal data redundancy and allows users to insert, delete, and update rows without causing data inconsistencies (anomalies)

Goal is to avoid anomalies

- **Insertion Anomaly** – adding new rows forces user to create duplicate data
- **Deletion Anomaly** – deleting rows may cause a loss of data that would be needed for other future rows
- **Modification Anomaly** – changing data in a row forces changes to other rows because of duplication

The **normalization process** involves decomposing relations with anomalies to produce smaller, well-structured relations

106

This table is NOT a relation

Violating the principle of “no multivalued attributes”. It’s as if OrderID 1006 has “ProductID” equaling 7, 5, and 4. You can’t have that in a true relation

Invoice data

| OrderID | Order Date | Customer ID | Customer Name | Customer Address | ProductID | Product Description | Product Finish | Product StandardPrice | Ordered Quantity |
|---------|------------|-------------|-------------------|------------------|-----------|----------------------|----------------|-----------------------|------------------|
| 1006 | 10/24/2015 | 2 | Value Furniture | Plano, TX | 7 | Dining Table | Natural Ash | 800.00 | 2 |
| | | | | | 5 | Writer's Desk | Cherry | 325.00 | 2 |
| | | | | | 4 | Entertainment Center | Natural Maple | 650.00 | 1 |
| 1007 | 10/25/2015 | 6 | Furniture Gallery | Boulder, CO | 11 | 4-Dr Dresser | Oak | 500.00 | 4 |
| | | | | | 4 | Entertainment Center | Natural Maple | 650.00 | 3 |

Issues

- OrderID blank in several rows.
- Don't explicitly see customer fields or order date for the 2nd, 3rd, and 5th rows. You might infer the values if this was a spreadsheet. This is not sufficient for a database.

107

Eliminating multivalued attributes creates a relation in First Normal Form (1NF)

INVOICE relation with unique rows and no multivalued attributes (in 1NF)

| OrderID | Order Date | Customer ID | Customer Name | Customer Address | ProductID | Product Description | Product Finish | Product StandardPrice | Ordered Quantity |
|---------|------------|-------------|-------------------|------------------|-----------|----------------------|----------------|-----------------------|------------------|
| 1006 | 10/24/2015 | 2 | Value Furniture | Plano, TX | 7 | Dining Table | Natural Ash | 800.00 | 2 |
| 1006 | 10/24/2015 | 2 | Value Furniture | Plano, TX | 5 | Writer's Desk | Cherry | 325.00 | 2 |
| 1006 | 10/24/2015 | 2 | Value Furniture | Plano, TX | 4 | Entertainment Center | Natural Maple | 650.00 | 1 |
| 1007 | 10/25/2015 | 6 | Furniture Gallery | Boulder, CO | 11 | 4-Dr Dresser | Oak | 500.00 | 4 |
| 1007 | 10/25/2015 | 6 | Furniture Gallery | Boulder, CO | 4 | Entertainment Center | Natural Maple | 650.00 | 3 |

Gaps filled

This is a relation, but not a well-structured one . . .

Note: the extensive data duplication (will cause anomalies)

All relations are in 1st Normal Form (1NF)

108

Relation in First Normal Form (1NF) can still suffer from anomalies*

INVOICE relation with unique rows and no multivalued attributes (in 1NF)

| OrderID | Order Date | Customer ID | Customer Name | Customer Address | ProductID | Product Description | Product Finish | Product StandardPrice | Ordered Quantity |
|---------|------------|-------------|-------------------|------------------|-----------|----------------------|----------------|-----------------------|------------------|
| 1006 | 10/24/2015 | 2 | Value Furniture | Plano, TX | 7 | Dining Table | Natural Ash | 800.00 | 2 |
| 1006 | 10/24/2015 | 2 | Value Furniture | Plano, TX | 5 | Writer's Desk | Cherry | 325.00 | 2 |
| 1006 | 10/24/2015 | 2 | Value Furniture | Plano, TX | 4 | Entertainment Center | Natural Maple | 650.00 | 1 |
| 1007 | 10/25/2015 | 6 | Furniture Gallery | Boulder, CO | 11 | 4-Dr Dresser | Oak | 500.00 | 4 |
| 1007 | 10/25/2015 | 6 | Furniture Gallery | Boulder, CO | 4 | Entertainment Center | Natural Maple | 650.00 | 3 |

Fields pertaining to Order Fields pertaining to Customer Fields pertaining to Product


These anomalies exist because there are multiple themes (entity types) in one relation... results in duplication and unnecessary dependency between the entities

* It can suffer from insertion, deletion, and update anomalies just like the EMPLOYEES2 relation we already saw

109

Functional Dependence exists when an attribute uniquely determines the value of other attributes

Dependency Diagram for EMPLOYEE1



An employee's name, department and salary are **functionally dependent** on his/her EmpID. That is, there can only be one name, department, and salary for each EmpID

Functional dependencies also written as follows


| Determinant | Dependents |
|--------------------------|--|
| EmpID → | Name, DeptName, Salary |
| VIN → | Make, Model, Colour |
| ISBN → | Title, FirstAuthorName, Publisher |
| EmpID, CourseID → | DateCompleted |

Finding functional dependencies is also a way of finding candidate PKs

110

Partial Functional Dependence is associated with data duplication and data anomalies

Dependency Diagram for EMPLOYEE2

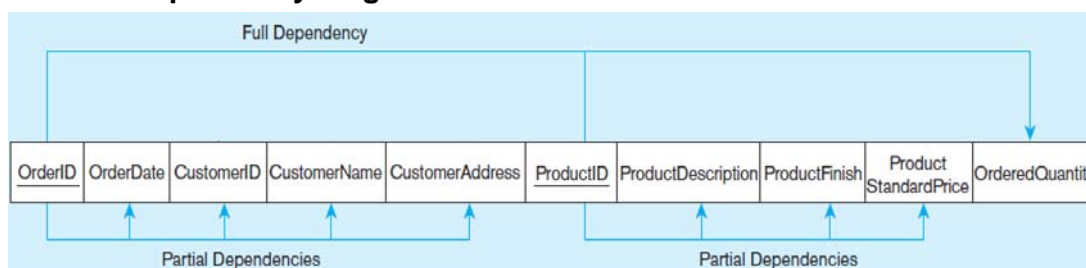


| Functional Dependency | Implication |
|---|--|
| EmpID, CourseTitle → DateCompleted | Only EmpID, CourseTitle is a candidate key and therefore must be the PK for this relation |
| EmpID → Name, DeptName, Salary | Three attributes are determined by only part of the PK for the whole relation This is called partial functional dependence |

111

Second Normal Form (2NF) is 1NF PLUS every non-key attribute **FULLY** functionally dependent on **ENTIRE** PK (i.e. no partial functional dependencies)

Functional dependency diagram for INVOICE



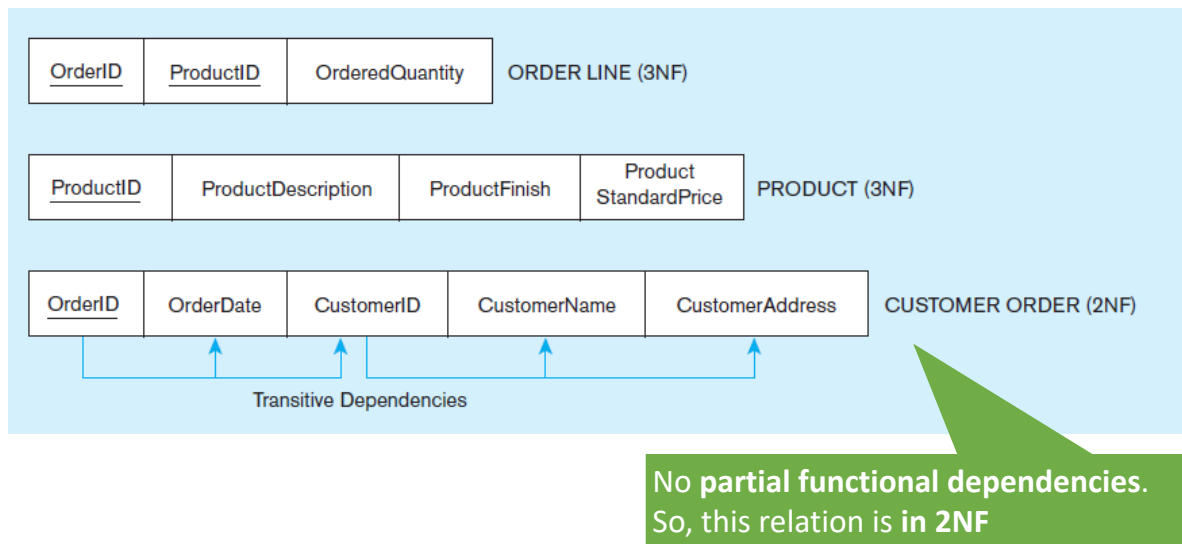
OrderID, ProductID → OrderQuantity

OrderID → OrderDate, CustomerID, CustomerName, CustomerAddress
ProductID → ProductDescription, ProductFinish, ProductStandardPrice

Two partial functional dependencies.
So, this relation is **NOT** in 2NF

112

Removing partial dependencies by splitting table into three (customer order, product, order line) gets us to 2NF




Partial dependencies are removed, but there are still transitive dependencies

113

Third Normal Form (3NF) is 2NF PLUS no transitive dependencies (functional dependencies on non-PK attributes)

Functional dependency diagram for CUSTOMER ORDER



OrderID → CustomerID → CustomerName
OrderID → CustomerID → CustomerAddress

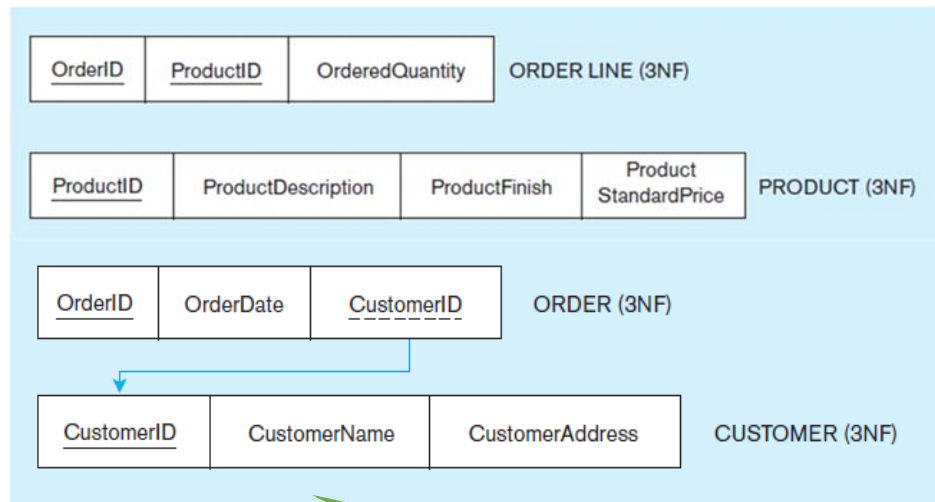
There are still transitive dependencies. So, this relation is not in 3NF

Called **transitive**, because PK is determinant for another attribute, which in turn is a determinant for a third attribute

Still have data redundancy since we would have to store customer name and address for each order

114

We reach Third Normal form (3NF)... by splitting out the customer attributes*



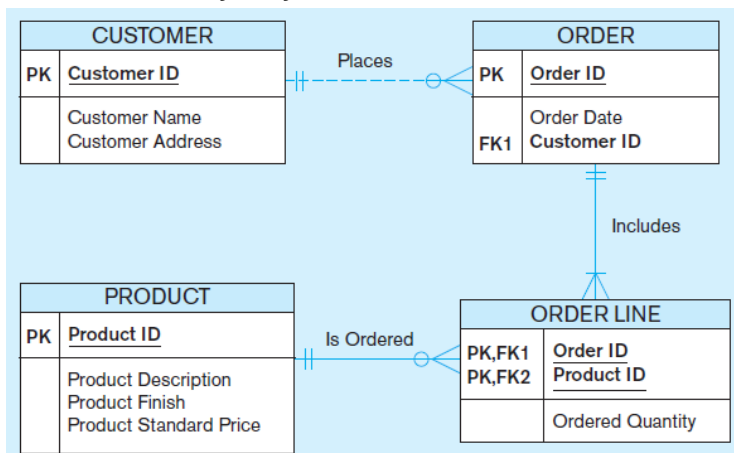
Transitive functional dependencies removed. So, this relation is in 3NF

* More generally, non-key determinant with transitive dependencies goes into a new table; non-key determinant becomes PK in new table and stays as FK in old table

115

Finally we reach Third Normal form (3NF)... now we have a well-structured database

Normalization, yielded four separate relations where there was initially only one



Original table

- Had three separate entities wrapped into single relation (customer, product, and order)
- Hid M:N relationship between orders and products

If we had started with a thorough ERD, we would not have wound up with one big table containing multiple themes



Illustrates danger of going straight to logical database design without performing conceptual model analysis.... but it happens often

116

In-class exercise: Normalization

15-minute
breakout

¶ The following is a sample of a spreadsheet used to track traffic violation at a college

| Parking Ticket Table | | | | | | | | | |
|----------------------|--------|--------|----------|--------|---------|----------|------------|------|-------|
| St ID | L Name | F Name | Phone No | St Lic | Lic No | Ticket # | Date | Code | Fine |
| 38249 | Brown | Thomas | 111-7804 | FL | BRY 123 | 15634 | 10/17/2018 | 2 | \$25 |
| | | | | | | 16017 | 11/13/2018 | 1 | \$15 |
| 82453 | Green | Sally | 391-1689 | AL | TRE 141 | 14987 | 10/05/2018 | 3 | \$100 |
| | | | | | | 16293 | 11/18/2018 | 1 | \$15 |
| | | | | | | 17892 | 12/13/2018 | 2 | \$25 |

¶ NOTE that

- Violation code 1 – expired traffic meter
- Violation code 2 – no parking permit
- Violation code 3 – accessible violation

¶ Develop a set of relations in 3NF

Use this format: ORDER OrderID OrderDate CustomerID

¶ Capture your answer on a PowerPoint page

¶ Submit your solution to <https://forms.gle/Kr9TjbQGhwwAGgwMA> One submission per breakout team

117

¶ Components of Relational Data Model

¶ Converting ERD into Database

¶ Database Normalization

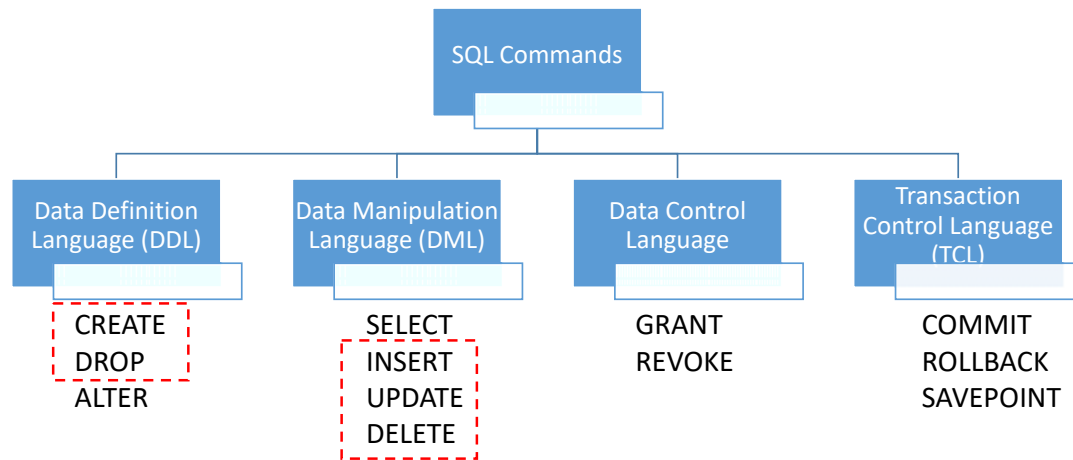
¶ Hands On

- SQL – Data Definition Language (DDL)
- SQL – INSERT, UPDATE and DELETE queries

¶ For next time

118

Basics of SQL DDL as well as INSERT, UPDATE, and DELETE queries covered in today's hands on scripts and videos



In computer programming, create, read, update, and delete (**CRUD**) are the four basic functions of persistent storage. This is how SQL DML queries map to CRUD

| CRUD | SQL |
|------------|--------|
| C - Create | Insert |
| R - Read | Select |
| U - Update | Update |
| D - Delete | Delete |

119

The script that creates the AP database (part 1)

```
-- create the database
DROP DATABASE IF EXISTS ap;
CREATE DATABASE ap;

-- select the database
USE ap;

-- create the tables
CREATE TABLE general_ledger_accounts
(
    account_number        INT          PRIMARY KEY,
    account_description    VARCHAR(50)  UNIQUE
);

CREATE TABLE terms
(
    terms_id              INT          PRIMARY KEY
                        AUTO_INCREMENT,
    terms_description      VARCHAR(50)  NOT NULL,
    terms_due_days         INT          NOT NULL
);
```

120

The script that creates the AP database (part 2)

```
CREATE TABLE vendors
(
    vendor_id            INT            PRIMARY KEY
                        AUTO_INCREMENT,
    vendor_name          VARCHAR(50)    NOT NULL
                        UNIQUE,
    vendor_address1      VARCHAR(50),
    vendor_address2      VARCHAR(50),
    vendor_city          VARCHAR(50)    NOT NULL,
    vendor_state         CHAR(2)        NOT NULL,
    vendor_zip_code      VARCHAR(20)    NOT NULL,
    vendor_phone         VARCHAR(50),
    vendor_contact_last_name VARCHAR(50),
    vendor_contact_first_name VARCHAR(50),
    default_terms_id     INT            NOT NULL,
    default_account_number INT          NOT NULL,
    CONSTRAINT vendors_fk_terms
        FOREIGN KEY (default_terms_id)
        REFERENCES terms (terms_id),
    CONSTRAINT vendors_fk_accounts
        FOREIGN KEY (default_account_number)
        REFERENCES general_ledger_accounts (account_number)
);
```

121

The script that creates the AP database (part 3)

```
CREATE TABLE invoices
(
    invoice_id          INT            PRIMARY KEY
                        AUTO_INCREMENT,
    vendor_id           INT            NOT NULL,
    invoice_number       VARCHAR(50)    NOT NULL,
    invoice_date         DATE           NOT NULL,
    invoice_total        DECIMAL(9,2)   NOT NULL,
    payment_total        DECIMAL(9,2)   NOT NULL    DEFAULT 0,
    credit_total         DECIMAL(9,2)   NOT NULL    DEFAULT 0,
    terms_id            INT            NOT NULL,
    invoice_due_date     DATE           NOT NULL,
    payment_date         DATE,
    CONSTRAINT invoices_fk_vendors
        FOREIGN KEY (vendor_id)
        REFERENCES vendors (vendor_id),
    CONSTRAINT invoices_fk_terms
        FOREIGN KEY (terms_id)
        REFERENCES terms (terms_id)
);
```

Variety of datatypes used in examples. See full details of MySQL datatypes here <https://dev.mysql.com/doc/refman/8.0/en/data-types.html>

- [11.1 Data Type Overview](#)
- [11.2 Numeric Types](#)
- [11.3 Date and Time Types](#)
- [11.4 String Types](#)
- [11.5 Spatial Data Types](#)
- [11.6 The JSON Data Type](#)

122

The script that creates the AP database (part 4)

```
CREATE TABLE invoice_line_items
(
    invoice_id            INT            NOT NULL,
    invoice_sequence      INT            NOT NULL,
    account_number        INT            NOT NULL,
    line_item_amount      DECIMAL(9,2)   NOT NULL,
    line_item_description  VARCHAR(100)   NOT NULL,
    CONSTRAINT line_items_pk
        PRIMARY KEY (invoice_id, invoice_sequence),
    CONSTRAINT line_items_fk_invoices
        FOREIGN KEY (invoice_id)
        REFERENCES invoices (invoice_id),
    CONSTRAINT line_items_fk_accounts
        FOREIGN KEY (account_number)
        REFERENCES general_ledger_accounts (account_number)
);

-- create an index
CREATE INDEX invoices_invoice_date_ix
ON invoices (invoice_date DESC);
```

123

Insert a single row without using a column list

```
INSERT INTO invoices VALUES
(115, 97, '456789', '2018-08-01', 8344.50, 0, 0, 1,
'2018-08-31', NULL)

(1 row affected)
```

Insert a single row using a column list

```
INSERT INTO invoices
    (vendor_id, invoice_number, invoice_total, terms_id,
    invoice_date, invoice_due_date)
VALUES
    (97, '456789', 8344.50, 1, '2018-08-01',
    '2018-08-31')

(1 row affected)
```

124

Update one column for multiple rows

```
UPDATE invoices  
SET terms_id = 1  
WHERE vendor_id = 95
```

(6 rows affected)

Update one column for one row

```
UPDATE invoices  
SET credit_total = credit_total + 100  
WHERE invoice_number = '97/522'
```

(1 row affected)

125

Delete one row

```
DELETE FROM general_ledger_accounts  
WHERE account_number = 306
```

(1 row affected)

Delete one row using a compound condition

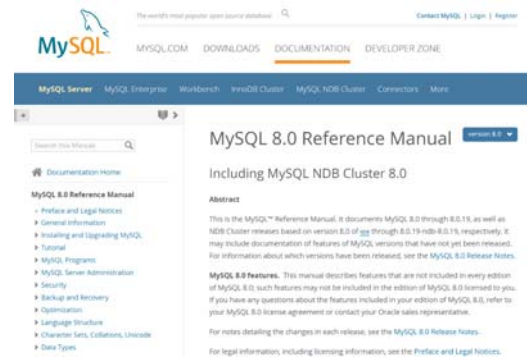
```
DELETE FROM invoice_line_items  
WHERE invoice_id = 78 AND invoice_sequence = 2
```

(1 row affected)

126

IMPORTANT

See 'hands on' scripts and videos for more examples... you will need those examples to complete the homework



¶ Impossible to cover all aspects of SQL in a short course... not even all the options available for the statements we do use. It is worthwhile getting comfortable using the documentation

- <https://dev.mysql.com/doc/>
- <https://dev.mysql.com/doc/refman/8.0/en/>

¶ Still having trouble with MySQL on your system? How many of you?

- Your root password for the MySQL server should be 'sesame80'
- See p46-47 in free pdf of chapter 2 of Murach's MySQL book for guide to starting MySQL server if it is not running
- LBS IT can provide support in getting you going (Philip Dainton-Smith?)

127

¶ Components of Relational Data Model

¶ Converting ERD into Database

¶ Database Normalization

¶ Hands On

¶ For next time

128

For next time

- ¶ Hands on activities (see Canvas page for session 2)
- ¶ Attend Thursday workshop (optional) if you think you might need more support before completing the homework
- ¶ Check out readings for this session and next
- ¶ Submit Homework Assignment #2 (team)
 - Covers more SQL queries and aspects of logical database design (details on Canvas)
 - Submit via Canvas by 10pm on Sunday

129

Appendix

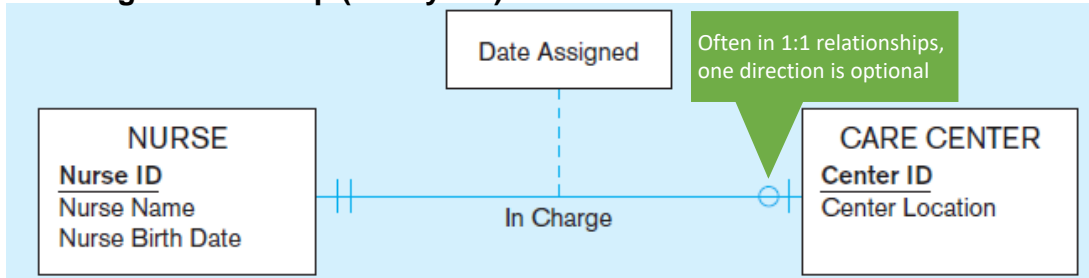
Additional Examples of Converting ERDs into Database

Hint: Pages on Supertype/subtype might help with homework assignment

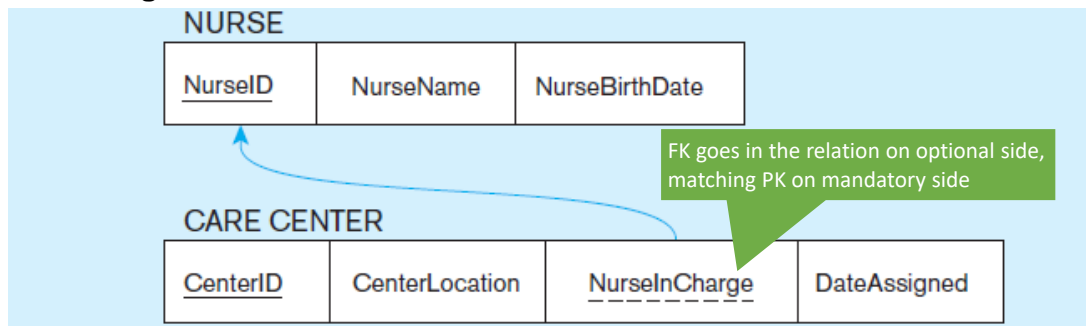
130

With a 1:1 relationship PK on mandatory side becomes a FK on optional side

In charge relationship (binary 1:1)



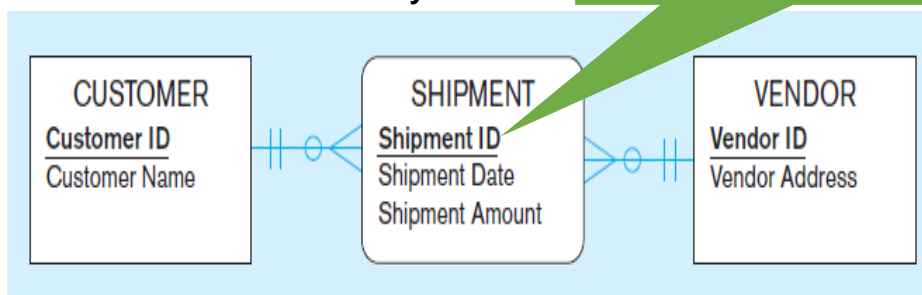
Resulting relations



131

Mapping associative entity with assigned identifier is slightly different (1/2)

SHIPMENT associative entity

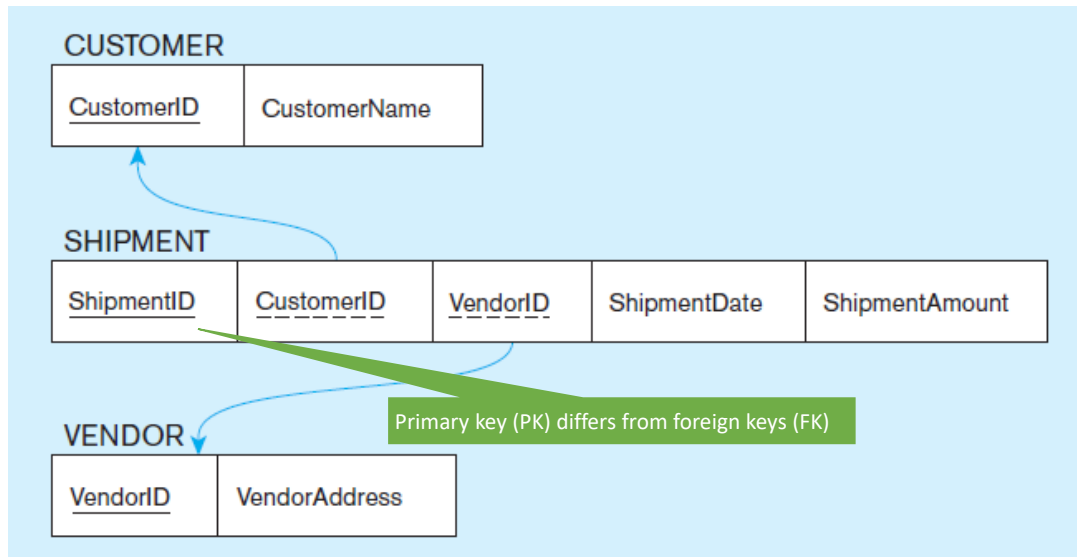


132

Mapping associative entity with assigned identifier is slightly different (2/2)

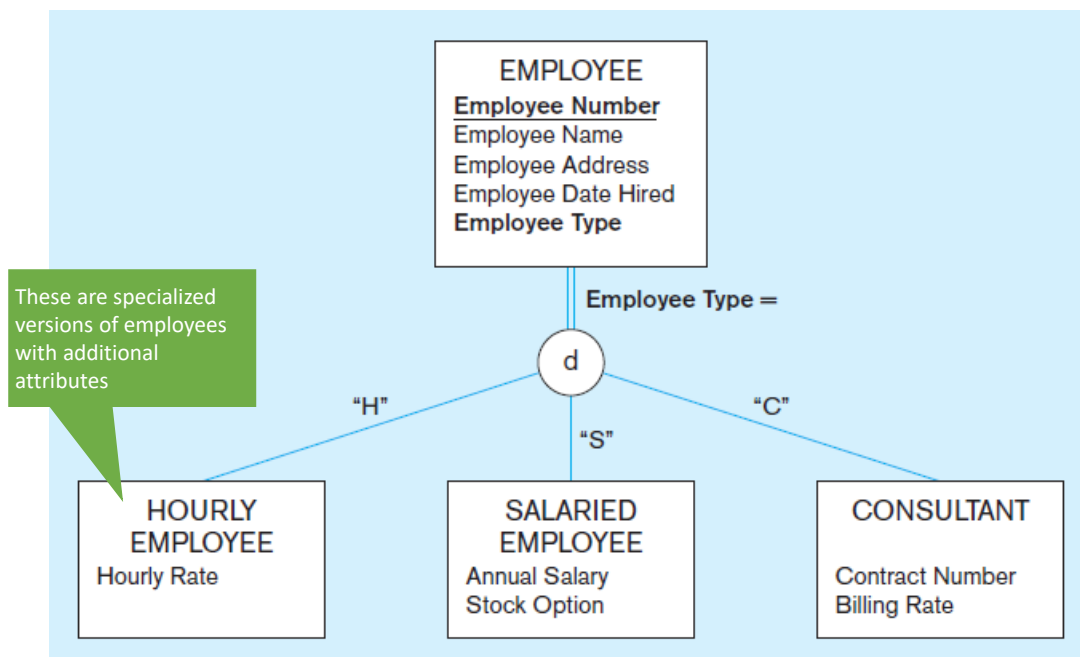


Three resulting relations

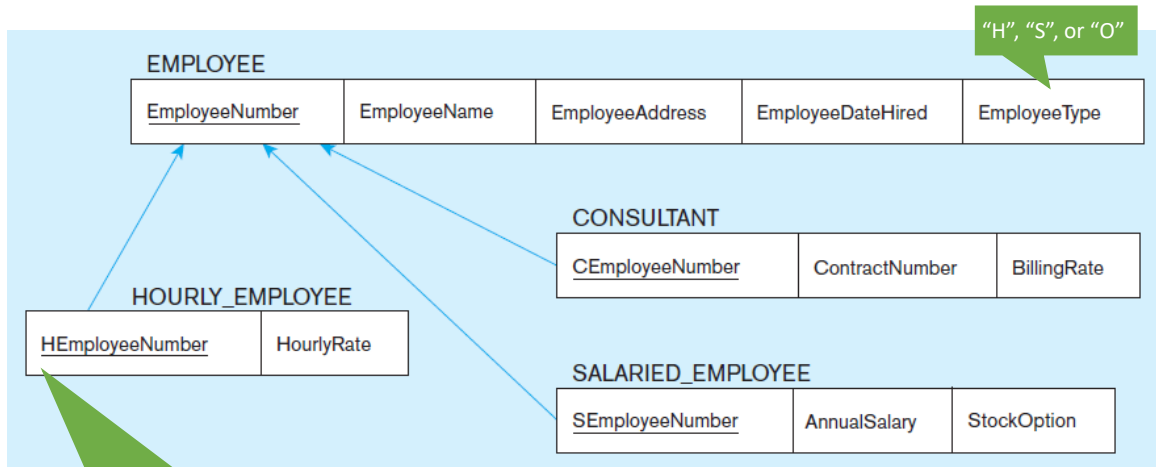


133

Supertype/subtype relationships can be modeled in ERDs



Supertype/subtypes are implemented in databases as one-to-one relationships



PK of the subtype relation is also a FK to the supertype relation.

If you see the entire PK is a FK in a database it implies a supertype-subtype relationship.